

AUDIENCE SIGNAL INTELLIGENCE — MONEY-BALLING CONTENT, PRODUCTIZED

Increase your hit rate with Audience Signal Intelligence.

Every year, billions in acquisition, licensing, and greenlight decisions are made on gut feel, comps, and star power. **HIT Predictor** decodes the narrative DNA of a title — conflict, stakes, warmth, hook, tone arc — and reads it against verified commercial performance across the largest certified content-identity spine in the industry. The output is an **Audience Signal Intelligence profile**: five forward-looking predictions per title, one instrument. That's not a comps deck. That's a trading terminal for content.

WHY NOW

Supply is finite. Production is unaffordable. Demand is winner-take-all. The margin for gut-feel greenlighting is gone.


735K MOVIES CATALOGUED GLOBALLY	360K TV SERIES & MINIS CATALOGUED GLOBALLY	-46% LA SHOOT DAYS SINCE 2022 (36.8K → 19.7K)	\$25B+ CONTENT WRITE-DOWNS ACROSS MAJOR STUDIOS	108B 1H'26 STREAMING MINUTES — CONCENTRATED IN JUST 5 TITLES	\$289B GLOBAL CONTENT MARKET UP FOR GRABS
------------------------------------	---	--	--	---	--

Sources: global title registry (Dec 2025) · FilmLA / ProdPro production tracking · national streaming measurement, 1H 2026 · SEC 10-K filings · #ShivSpeaks: "The Merger Paradox" & "Media Freakanomics: The Genre Arbitrage" (2026)

The Audience Signal Intelligence profile

One title in, five predictions out — each modeled against the ground truth already fused on the spine: **theatrical economics, audience viewership, ad delivery, streaming demand, and social conversation**, all resolved to a single certified identity per title.


- Box Office** Theatrical revenue & budget-multiplier band, from content-DNA and talent-value economics
- Licensing** Downstream library & catalog value — window-by-window revenue potential
- Linear Ad** Ad-supported delivery & monetizable audience across national telecast windows
- SVOD** Streaming viewership demand & retention contribution by platform profile
- Social** Engagement velocity, share of voice & buzz durability across social platforms



ILLUSTRATIVE SIGNAL PROFILE · SAMPLE TITLE · 0-100 SIGNAL INDEX PER SURFACE

How it works

The hard problem was never the model — it's the join. Connecting a title's creative attributes to its true commercial performance across fragmented ID systems is where every prior attempt died. HIT Predictor's **12.96M-row certified identity spine** solves it — a certified 220K-title candidate pool, with title-integrity guards that caught and rejected a 25.6% bad-match rate that would silently poison anyone else's model.



The spine is the moat. Anyone can prompt an LLM about a script. Only HIT Predictor can test the answer against what audiences — and markets — actually did.

WHERE IT STANDS — AND WHY WE PUBLISH IT

The v0.1 pilot found real signal in summary text — and proved it isn't enough.

Nearly 10,000 titles scored across two runs on the linear-audience surface. Marketing-copy proxies produce **replicable but small** effects (<2% of variance) — a finding most vendors would bury. We publish it, because it defines the build: the levers are **script-level text, per-demographic targets, and time-series windows** — all on the spine already. Candor is the credibility. The same discipline — title-integrity guards, certified crosswalks, dedup audits — is what makes every signal trustable when it lands.

Why this team

The concept has shipped before.

The founder architected trait-based movie-performance prediction inside a top-5 US media company — deployed against a live daytime acquisition slate — and built audience forecasting engines running at 95%+ daily accuracy in production.

Award-winning methodology.

INFORMS Practice Award and FICO Decisions Award for ML and optimization systems generating nine years of measured revenue lift on national television.

Clean, owned IP.

The Parent Key identity spine and matching methodology are first-time commercial implementations — no legacy licenses, no encumbrances.

Roadmap to v1

- NOW

Scale on the promoted pool. Re-run scoring against the certified 220K-title candidate set; light up box-office and social surfaces already fused on the spine.
- NEXT

Script-level text. Replace summary proxies with dialogue and screenplay features — where the narrative signal actually lives.
- THEN

Per-demo & per-window models. Segment-level audience targets and window-by-window licensing value, not blended totals that average the signal away.
- V1

The Audience Signal Intelligence instrument. Five-signal slate-scoring API + bankability index inside the terminal.

Who it's for

Content creators — networks, streamers, and studios pressure-testing greenlights and pricing library acquisitions. **Distributors** valuing catalogs window by window. Anywhere a seven-figure content decision currently rests on a comps slide.

Seeking 4 individuals to co-develop the Audience Signal Intelligence platform.

Shiv Ashish Sehgal

Founder

Former CPO, RSG Media · UCLA-NUS EMBA

What the instrument shows you.


Five working views from the Audience Signal Intelligence terminal — how a slate, a deal, a genre strategy, and a single title's revenue stack look when the gut-feel is replaced with signal. Real spine data where marked; illustrative where marked.

The Slate View — signal on real titles

BUDGET · BOX OFFICE · MULTIPLE (PUBLIC FIGURES) · SIGNAL PROFILE 0-100

Six recent theatricals read through the instrument. The economics are real; the lesson is the moneyball one — **budget doesn't predict the multiple, signal does**. A \$100M animated title returned 13.6× while a \$275M tentpole returned 0.7×. Cost is what you pay; signal is what you get.

Box Office Licensing Linear Ad SVOD Social




Budget and worldwide box office per The Numbers / Box Office Mojo (approx.; Avatar WoW budget estimates range \$350-400M). Poster imagery via TMDB. Signal profiles and composite badges are illustrative back-scores, not live model output — the demonstration is the read, not the prediction.

The Deal Quadrant

Every candidate plotted on what it costs against what the signal says it's worth. The upper-left is where deals are won.


SIGNAL VS. PRICE



The Genre Arbitrage

AUDIENCE DEMAND VS. NEW COMMISSIONS

The industry herds into drama and reality while comedy and animation audiences go underserved — the blind spot the signal exposes, title by title. **The real spine read backs it (n=206,766 series):** comedy's share of new series nearly halved 2010→2025 while drama's more than doubled — and drama converts to a second season less often.



ρ = 0.13

CATALOGUE POPULARITY VS MEASURED AUDIENCE — UNINFORMATIVE. THE INDUSTRY'S FAVORITE METRIC DOESN'T PREDICT VIEWING.

~8×

PRE-1990 SCRIPTED LIBRARY VS 2010+ ORIGINALS, MEDIAN AUDIENCE PER TELECAST (SUMMER WINDOW)

82.7%

OF ALL VIEWING MINUTES COME FROM THE HEAVIEST ONE-FIFTH OF VIEWERS

Demand/commission shares are directional, per published #ShivSpeaks analysis; the comedy/drama supply shift and the three reads above are measured on the platform's identity spine (2026 windows; unit-free ratios and shares; direction-level claims; genre tags are multi-label). Quadrant remains illustrative.

The Feature DNA

Live output from the n=43 scored pilot — real, not illustrative. Every title is decoded into eight narrative features; this is how they relate to each other. Conflict and darkness travel together (ρ 0.94); warmth cuts against both; humor stands alone.

This matrix is the QC layer of the instrument. Before a feature is trusted in a signal model, we measure what it actually captures — including the honest finding that features extracted from the same text co-move, and that none alone predicts delivery yet. Knowing that is the difference between a model and a vibe.


FEATURE-VS-FEATURE SPEARMAN P · 43 SCORED TITLES · TEAL = CO-MOVE · CORAL = OPPOSE

	HOOK	CONFLICT	EM. RNG	WARMTH	DARKNESS	HUMOR	SERIAL.	STAKES
HOOK	1.00	0.91	0.65	-0.46	0.89	-0.18	0.62	0.84
CONFLICT	0.91	1.00	0.65	-0.53	0.94	-0.19	0.65	0.90
EM. RNG	0.65	0.65	1.00	0.00	0.62	0.08	0.69	0.47
WARMTH	-0.46	-0.53	0.00	1.00	-0.61	0.46	-0.07	-0.58
DARKNESS	0.89	0.94	0.62	-0.61	1.00	-0.21	0.57	0.86
HUMOR	-0.18	-0.19	0.08	0.46	-0.21	1.00	0.04	-0.45
SERIAL.	0.62	0.65	0.69	-0.07	0.57	0.04	1.00	0.52
STAKES	0.84	0.90	0.47	-0.58	0.86	-0.45	0.52	1.00

The Revenue Stack — one title, every window

CUMULATIVE PREDICTED VALUE, \$M · SAMPLE TITLE

The five-signal profile translated into a window-by-window P&L — the number a distributor actually negotiates against. Signal colors carry through: each window is priced by the surface that drives it.



All figures on this page are illustrative sample outputs, not live model predictions — surfaces reach modeled ground truth on the roadmap sequence shown on page 1. Genre demand/commission shares are directional, per published #ShivSpeaks analysis.